

# Basic transformer architecture

AP & MF



**Preliminaries**

# Tokens, vocabulary, language models

- ▶  $\mathcal{V} = \{t_1, \dots, t_{n_{\mathcal{V}}}\}$  is the *vocabulary*
  - finite set of *tokens* of cardinality  $n_{\mathcal{V}}$
  - vocabulary is ordered by the bijective *indexing function*  $\text{Ind}: \mathcal{V} \rightarrow \{1, \dots, n_{\mathcal{V}}\}$ 
    - $\text{Ind}(t)$  is the (token) *index* of  $t$
- ▶ A *token sequence* of length  $n \in \mathbb{N}$  is a vector  $\mathbf{t} \in \mathcal{V}^n$  of  $n$  tokens from the vocabulary.
- ▶ The set of all finite token sequences forms the *language*  $\mathcal{L} = \{\mathbf{t} \in \mathcal{V}^n \mid n \in \mathbb{N}\}$ .
- ▶ A (*generative / decoder*) *language model* is a function  $\mathcal{M}_{\theta}: X \rightarrow \Delta(\mathcal{L})$ , parameterized on a set of model parameter values  $\theta \in \Theta$ , which maps some input set  $X$  onto a probability distribution over all (finite) sequences of tokens.

# Autoregressive vs. masked LMs


- ▶ An *autoregressive (next-token prediction) language model* is a function  $\mathcal{M}_\theta: \mathcal{Z} \rightarrow \Delta(\mathcal{V})$  which maps a token sequence onto a *next-token probability distribution*.
  - We write  $P_{\mathcal{M}}(t \mid \mathbf{t}; \theta)$  or, omitting the model's parameters,  $P_{\mathcal{M}}(t \mid \mathbf{t})$  for the model's prediction of the next-token probability for token  $t$ .
- ▶ A *masked (missing-token prediction) language model* is a function that predicts selected tokens inside of a given sequence of tokens.
  - The next-token prediction LMs can be considered a special case of the missing-token language models.



**Transformer architecture**

# General transformer architecture

- ▶ The transformer model processes an input token sequence in parallel across all positions using a chain —stack— of  $n_L$  *transformer blocks*.
  - We refer to the  $\ell^{\text{th}}$  element in the chain as the  $\ell^{\text{th}}$  *layer*.



# Embeddings, model size, transformer block

- ▶ An embedding  $\mathbf{x} \in \mathbb{R}^{d_{\mathcal{M}}}$  is an  $d_{\mathcal{M}}$ -dimensional vector.
  - $d_{\mathcal{M}}$  the *dimensionality of the embedding space*, aka. the *model size*
- ▶ The transformer block  $T^{(\ell)}$  for the  $\ell^{\text{th}}$  layer is a function  $T^{(\ell)}: \{1, \dots, n_c\} \times (\mathbb{R}^{d_{\mathcal{M}}})^{n_c} \rightarrow \mathbb{R}^{d_{\mathcal{M}}}$  which maps a the  $i^{\text{th}}$  input position of context  $C$  ( $1 \leq i \leq n_c$ ) and a sequence of context embeddings to a single embedding as output, such that:

$$T^{(\ell)}: \left\langle \underbrace{i}_{\text{index of token}}, \underbrace{\langle \mathbf{x}_1^{(\ell-1)}, \dots, \mathbf{x}_{n_c}^{(\ell-1)} \rangle}_{\text{embeddings at previous layer}} \right\rangle \mapsto \underbrace{\mathbf{x}_i^{(\ell)}}_{\text{embedding of } i \text{ at current layer}}$$

- For  $\ell = 1$ , input zero-layer token embeddings defined next.

# Zero-layer, initial token embeddings

- ▶ The zero-layer, initial token embedding,  $\mathbf{x}_i^{(0)} \in \mathbb{R}^{d_{\mathcal{M}}}$ , of input token  $i$  is computed as:

$$\mathbf{x}_i^{(0)} = F \left( \mathbf{W}_{E_{\text{Ind}(i)}}, \text{PosEmbed}(i) \right), \text{ where}$$

- $\mathbf{W}_E \in \mathbb{R}^{n_v, d_{\mathcal{M}}}$  is the *embedding matrix*
  - the values of  $\mathbf{W}_E$  are part of the model's trainable parameters
- $\text{PosEmbed}(i) \in \mathbb{R}^{d_{\mathcal{M}}}$  is a *positional embedding function*
- $F: \mathbb{R}^{d_{\mathcal{M}}} \times \mathbb{R}^{d_{\mathcal{M}}} \rightarrow \mathbb{R}^{d_{\mathcal{M}}}$  implements some way of combining the information from the (non-positional, fixed) embedding matrix with the positional information
  - $F$  can be as simple as mere addition



# Unembedding and token predictions

- ▶ Model predictions for position  $i$  are derived by applying the model's *unembedding matrix*  $U \in \mathbb{R}^{n_v, d_M}$  to  $\mathbf{x}_i^{(n_L)}$  (final layer embedding at position  $i$ ) to obtain a vector of *output logits*  $\text{logits}_i(\mathbf{t}) \in \mathbb{R}^{n_v}$ :

$$\text{logits}_i(\mathbf{t}) = U \mathbf{x}_i^{(n_L)}$$

- ▶ Logits are transformed into probabilities using softmax:


$$P_{\mathcal{M}}(t \mid \mathbf{t}, i) = \frac{\exp(\text{logits}_i(\mathbf{t})_{\text{Ind}(t)})}{\sum_{j=1}^{n_v} \exp(\text{logits}_i(\mathbf{t})_j)}$$

- (autoregressive) next-token prediction models predict the *next* token at  $i + 1$
- (masked) missing-token prediction models (may) predict the *current* token at  $i$



## **Anatomy of a transformer block**

# Anatomy of a transformer block



# Anatomy of a transformer block

- ▶ a transformer for layer  $\ell$  at position  $i$  is a function  $T^{(\ell)}: \langle i, \langle \mathbf{x}_1^{(\ell-1)}, \dots, \mathbf{x}_{n_c}^{(\ell-1)} \rangle \rangle \mapsto \mathbf{x}_i^{(\ell)}$  that maps a token index and a sequence of embeddings to a new embedding.

$$\mathbf{x}_i^{(\ell-1,*)} = \text{LayerNorm}^{(\ell,1)} \left( \mathbf{x}_i^{(\ell-1)} + \text{Attention}_i^{(\ell)} \left( \mathbf{x}_1^{(\ell-1)}, \dots, \mathbf{x}_{n_c}^{(\ell-1)} \right) \right)$$

$$\mathbf{x}_i^{(\ell)} = \text{LayerNorm}^{(\ell,2)} \left( \mathbf{x}_i^{(\ell-1,*)} + \text{FFN}^{(\ell)} \left( \mathbf{x}_i^{(\ell-1,*)} \right) \right)$$

- adding a *residual connection* to an operation or function  $F$  applied to input  $x$ , is to compute the (parameter-free) function  $G(x): x \mapsto x + F(x)$

# Layer normalization

- ▶ The *layer-normalization function*  $\text{LayerNorm}: \mathbb{R}^{d_{\mathcal{M}}} \rightarrow \mathbb{R}^{d_{\mathcal{M}}}$  is included for training efficiency and defined as:

$$\text{LayerNorm}(\mathbf{x})^{(\ell,k)} = \gamma^{(\ell,k)} \odot \frac{\mathbf{x} - \mu(\mathbf{x})}{\sigma(\mathbf{x})} + \beta^{(\ell,k)},$$

where:

- $\mu(\mathbf{x}) = \frac{1}{d} \sum_{i=1}^{d_{\mathcal{M}}} x_i$  is the mean of the values in the embedding,
- $\sigma(\mathbf{x}) = \sqrt{\frac{1}{d} \sum_{i=1}^{d_{\mathcal{M}}} (x_i - \mu(\mathbf{x}))^2 + \epsilon}$  is the standard deviation with a small constant  $\epsilon > 0$  for numerical stability,
- $\gamma^{(\ell,k)}, \beta^{(\ell,k)} \in \mathbb{R}^{d_{\mathcal{M}}}$  are learned parameters (scale and shift), one for each layer  $\ell$  and occurrence  $k$  of the layer-normalization operation in the transformer block, and
- $\odot$  denotes elementwise multiplication.

# Feed-forward network

- ▶ The *feed-forward network*  $\text{FFN}^{(\ell)} : \mathbb{R}^{d_{\mathcal{M}}} \rightarrow \mathbb{R}^{d_{\mathcal{M}}}$  usually has one hidden layer of size  $d_{\text{FFN}}$  ( $\approx 2d_{\mathcal{M}}$ ):

$$\text{FFN}^{(\ell)} \left( \mathbf{x}_i^{(\ell,*)} \right) = \mathbf{W}_{\text{out}}^{(\ell)} \left( \phi \left( \mathbf{W}_{\text{in}}^{(\ell)} \mathbf{x}_i^{(\ell,*)} \right) + \boldsymbol{\beta}^{(\text{in},\ell)} \right) + \boldsymbol{\beta}^{(\text{out},\ell)},$$

where

- $\mathbf{W}_{\text{in}}^{(\ell)} \in \mathbb{R}^{d_{\text{FFN}} \times d_{\mathcal{M}}}$  is the mapping from the input layer to the hidden layer
- $\mathbf{W}_{\text{out}}^{(\ell)} \in \mathbb{R}^{d_{\mathcal{M}} \times d_{\text{FFN}}}$  is the mapping from hidden layer to output layer,
- $\boldsymbol{\beta}^{(\text{in},\ell)}$  and  $\boldsymbol{\beta}^{(\text{out},\ell)}$  are bias vectors, and
- $\phi$  is some (non-linear) activation function.

# Attention module

- ▶ The *attention module* maps a position index and a sequence of embeddings to another embedding

$$\text{Attention}^{(\ell)}: \{1, \dots, n_C\} \times \left(\mathbb{R}^{d_M}\right)^{n_C} \rightarrow \mathbb{R}^{d_M}$$

such that:

$$\text{Attention}^{(\ell)}\left(i, \mathbf{x}_1^{(\ell-1)}, \dots, \mathbf{x}_{n_C}^{(\ell-1)}\right) = \mathbf{W}_O \mathbf{z}, \text{ with}$$

$$\mathbf{z} = \underbrace{A\mathbf{H}^{\ell,1}\left(i, \mathbf{x}_1^{(\ell-1)}, \dots, \mathbf{x}_{n_C}^{(\ell-1)}\right) \parallel \dots \parallel A\mathbf{H}^{\ell,n_H}\left(i, \mathbf{x}_1^{(\ell-1)}, \dots, \mathbf{x}_{n_C}^{(\ell-1)}\right)}_{\text{concatenation of outputs of } n_H \text{ attention heads}}$$

where  $\mathbf{W}_O \in \mathbb{R}^{d_M \times n_H d_h}$  is an output mapping, where  $d_h$  is the dimension of the output of a single attention head.

# Attention head

- ▶ The *attention head*  $h$  for is a function  $\text{AH}^{\ell,h}: \{1, \dots, n_c\} \times (\mathbb{R}^{d_M})^{n_c} \rightarrow \mathbb{R}^{d_h}$ :

$$\text{AH}^{\ell,h} \left( i, \mathbf{x}_1^{(\ell-1)}, \dots, \mathbf{x}_{n_c}^{(\ell-1)} \right) = \sum_{j=1}^{n_c} \text{Weight}^{(\ell,h)} \left( \mathbf{x}_i^{(\ell-1)}, \mathbf{x}_j^{(\ell-1)} \right) \text{Value}^{(\ell,h)} \left( \mathbf{x}_j^{(\ell-1)} \right)$$

- ▶ The *value function*  $\text{Value}: \mathbb{R}^{d_M} \rightarrow \mathbb{R}^{d_V}$  maps an embedding of size  $d_M$  to a value vector of size  $d_V$ :

$$\text{Value}^{(\ell,h)}(\mathbf{x}) = \mathbf{W}_V^{(\ell,h)} \mathbf{x}$$

- ▶ The *attention weight function* is defined as a softmax over numerical scores (where we assume that  $\exp(-\infty) = 0$  and  $d_k$  is the size of key / query vectors):

$$\text{Weight}^{(\ell,h)}(\mathbf{x}_i, \mathbf{x}_j) = \frac{\exp \left( d_k^{-1/2} \text{Score}^{(\ell,h)}(\mathbf{x}_i, \mathbf{x}_j) \right)}{\sum_{j'=1}^C \exp \left( d_k^{-1/2} \text{Score}^{(\ell,h)}(\mathbf{x}_i, \mathbf{x}_{j'}) \right)}$$



# Key, query, scores, masking

- ▶ The *key and query functions* map embeddings from  $\mathbb{R}^M$  to vectors of type  $\mathbb{R}^{d_k}$  for some dimensionality  $d_k$ :

$$\text{Key}^{(\ell, h)}(\mathbf{x}) = \mathbf{W}_K^{(\ell, h)} \mathbf{x}$$

$$\text{Query}^{(\ell, h)}(\mathbf{x}) = \mathbf{W}_Q^{(\ell, h)} \mathbf{x}$$

- ▶ The *score function* maps to embeddings from  $\mathbb{R}^M$  onto a non-normalized score, taking information from non-masked tokens into account:

$$\text{Score}^{(\ell, h)}(\mathbf{x}_i, \mathbf{x}_j) = \begin{cases} -\infty & \text{if token } j \text{ is masked for } i \\ \text{Query}^{(\ell, h)}(\mathbf{x}_i) \cdot \text{Key}^{(\ell, h)}(\mathbf{x}_j) & \text{otherwise.} \end{cases}$$